

LinksIn

Amy Tu
Junhong Wang
Omar Tleimat
Cheuk Yin Phipson Lee



- 1. Overview**
- 2. Load Testing**
- 3. Future Development**
- 4. Conclusion**

Project Overview

Demo project, application architecture

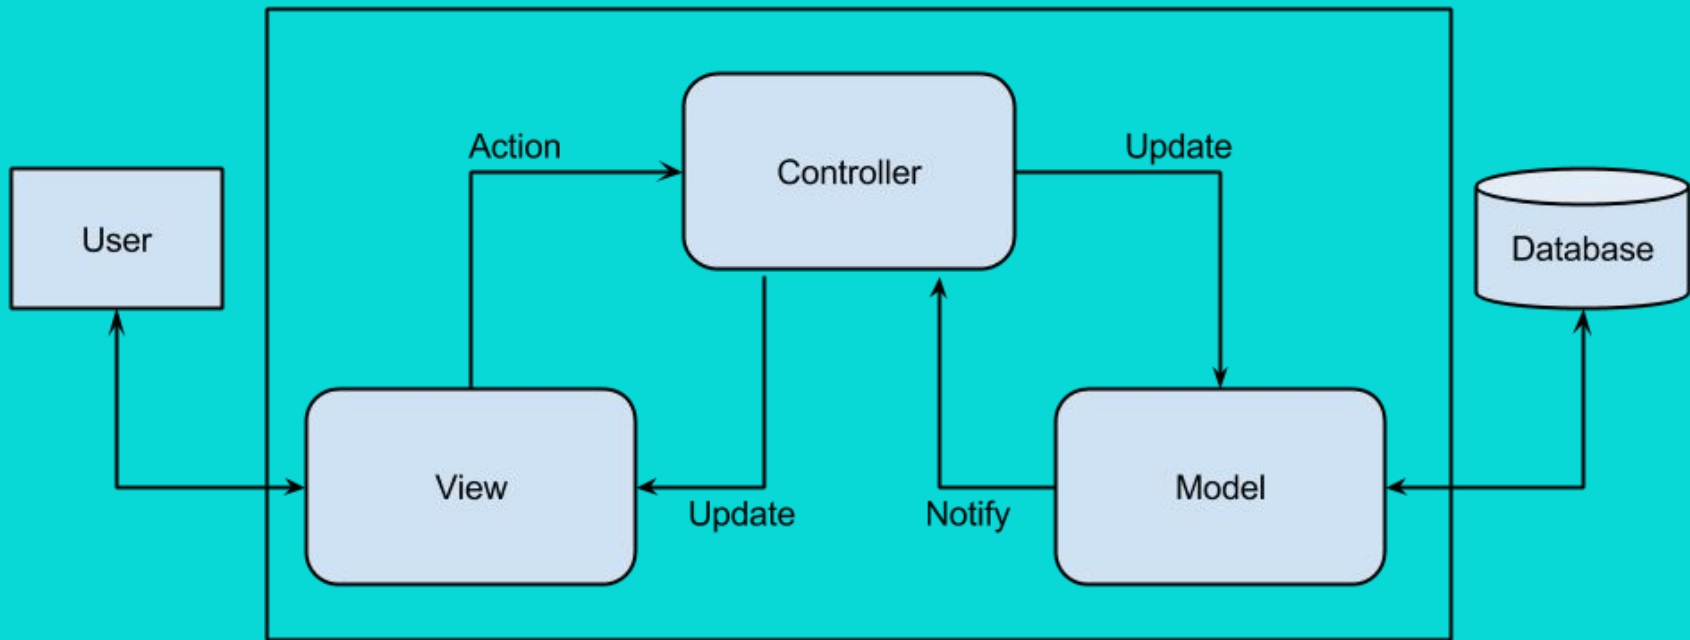




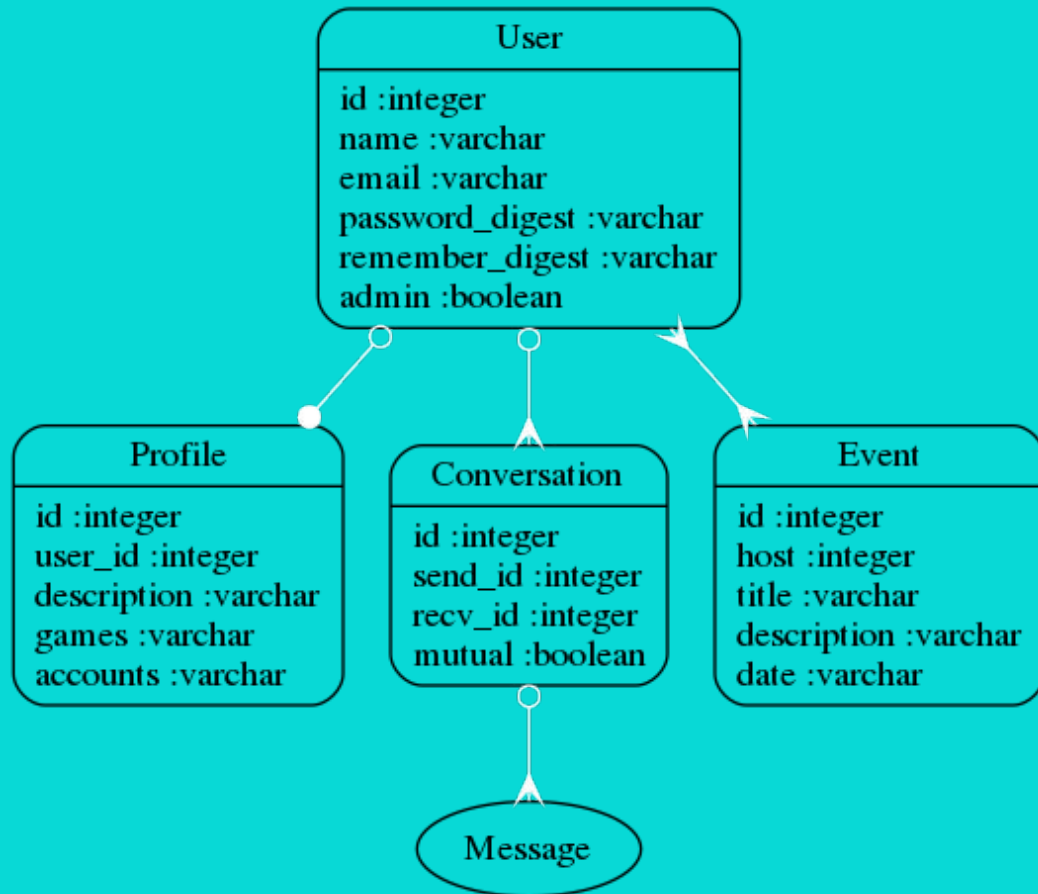
Application Architecture

Model-View-Controller (MVC)

Model-View-Controller Architecture

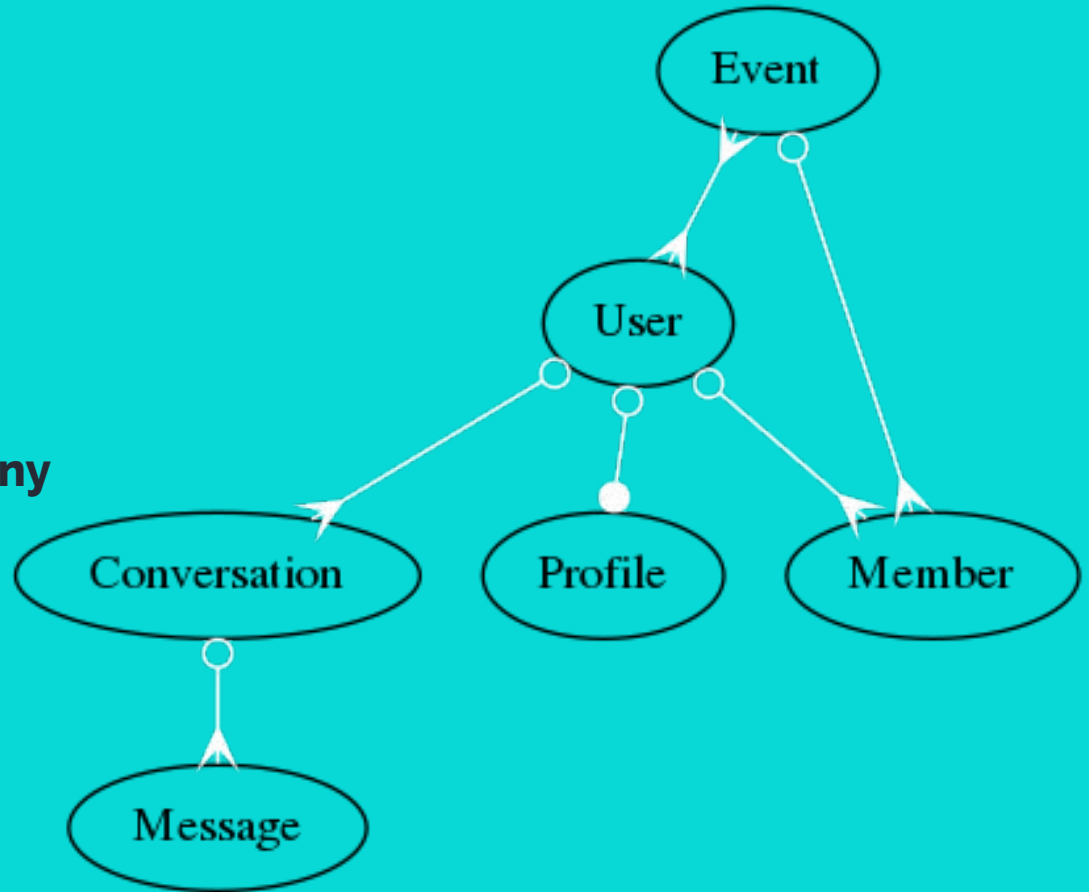


UML Diagram




Design Decisions

has_many, through =>
vs
has_and_belongs_to_many



Experiment & Results

Critical path, optimization before and afters, future optimizations



Test Phases

Test Phases

Phases	Users/sec
1	1
2	1.5
3	2
4	4
5	6
6	10
7	16
8	20
9	25
10	35
11	45
12	55




Critical Path

Critical Path

Step	Operation
1	Go to home page
2	Go to login page
3	Fill out login form
4	Log in and redirected to home page
5	Swipe right on user 1 and redirect to messages
6	Messages to user 1
7	Go to home page
8	Repeat step 5 - 7 for user 2, 3, 4, and 5
9	Go to events page
10	Go to event creation page
11	Create an event


Vertical Scaling



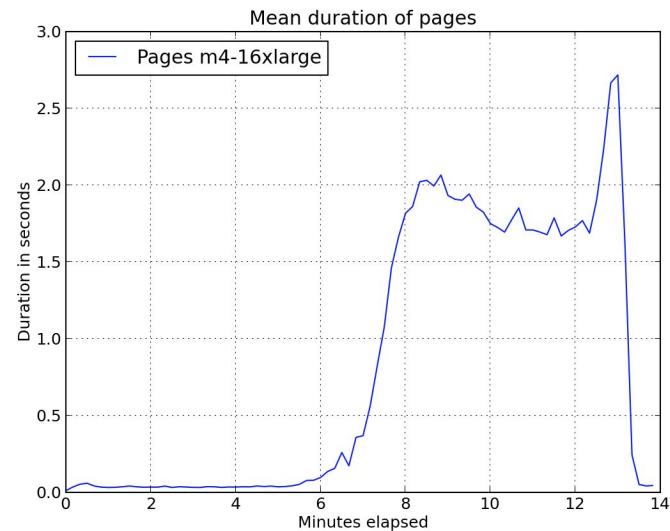
Vertical Scaling

Instance	vCPU	Memory (GiB)	Price
t3.micro	2	1	\$7.4/month
t3.large	2	8	\$60.0/month
t3.xlarge	4	16	\$119.8/month
t3.2xlarge	8	32	\$239.6/month
m4.4xlarge	16	64	\$576/month
m4.16xlarge	64	256	\$2304/month

Vertical Scaling



t3-micro




m4-16xlarge

Vertical Scaling

Instance	Price	Max Phase Handled
t3.micro	\$7.4/month	3 (2 users/sec)
t3.large	\$60.0/month	3 (2 users/sec)
t3.xlarge	\$119.8/month	5 (6 users/sec)
t3.2xlarge	\$239.6/month	5 (6 users/sec)
m4.4xlarge	\$576/month	5 (6 users/sec)
m4.16xlarge	\$2304/month	6 (10 users/sec)


Horizontal Scaling




Horizontal Scaling

Instance	vCPU	Memory (GiB)	Price
t3.micro x 1	2 x 1	1 x 1	\$7.49/month
t3.micro x 2	2 x 2	1 x 2	\$14.98/month
t3.micro x 4	2 x 4	1 x 4	\$30.00/month
t3.micro x 8	2 x 8	1 x 8	\$59.90/month
t3.micro x 16	2 x 16	1 x 16	\$119.81/month
t3.micro x 32	2 x 32	1 x 32	\$239.62/month

Horizontal Scaling



t3-micro x 1



t3-micro x 32

Horizontal Scaling


Instance	Price	Max Phase Handled
t3.micro x 1	\$7.49/month	3 (2 users/sec)
t3.micro x 2	\$14.98/month	5 (6 users/sec)
t3.micro x 4	\$30.00/month	5 (6 users/sec)
t3.micro x 8	\$59.90/month	5 (6 users/sec)
t3.micro x 16	\$119.81/month	5 (6 users/sec)
t3.micro x 32	\$239.62/month	6 (10 users/sec)



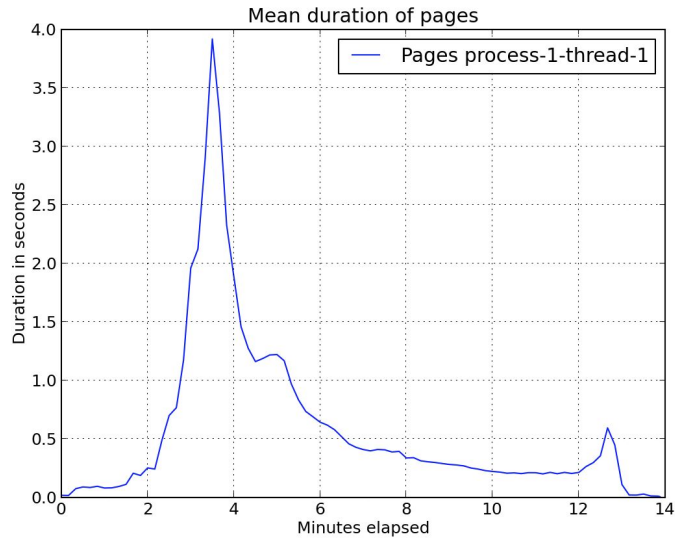
Horizontal Scaling

Save \$2064.38/month

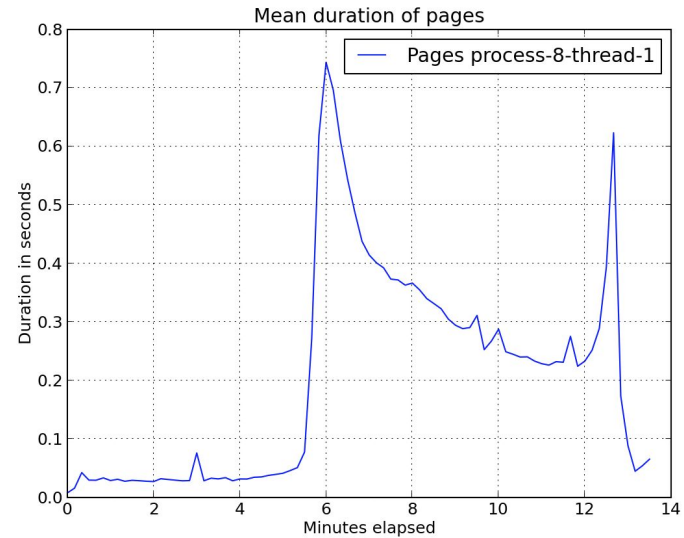
Processes/Threads



Processes/Threads




process x 1
thread x 1




process x 8
thread x 1

Processes/Threads



process x 1
thread x 1

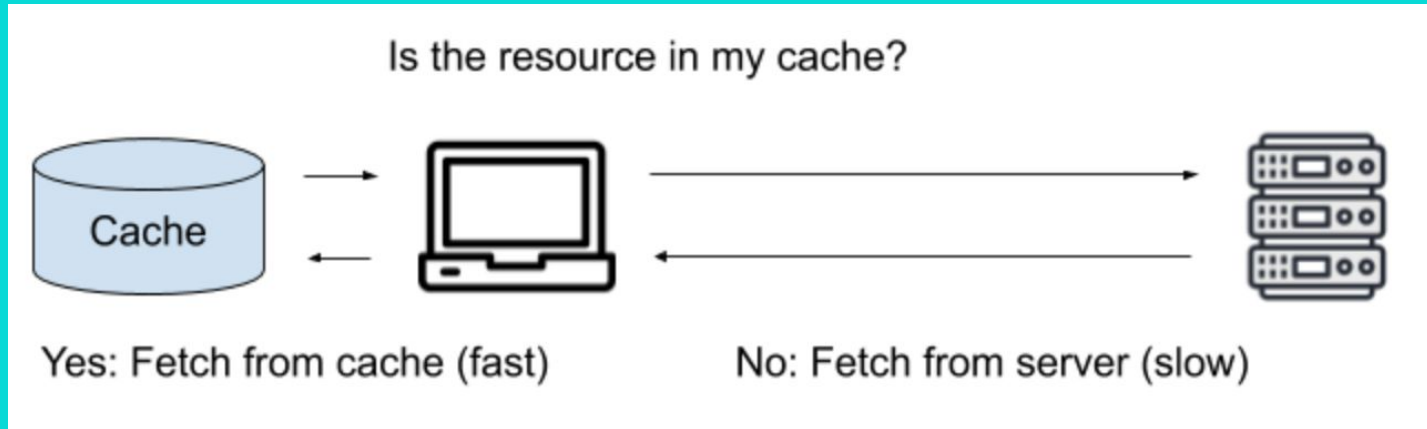


process x 1
thread x 8

Processes/Threads

Instance	Max Phase Handled
process x 1 & thread x 1	1 (1 user/sec)
process x 8 & thread x 1	5 (6 users/sec)
process x 1 & thread x 8	3 (2 users/sec)

Client-side Caching



Client-side caching


user_controller.rb

```
if stale?([@user, @user.updated_at, @user.events])  
  @user = User.find(params[:id])  
end
```


event_controller.rb

```
if stale?([@event, @event.updated_at])  
  @event = Event.find(params[:id])  
end
```

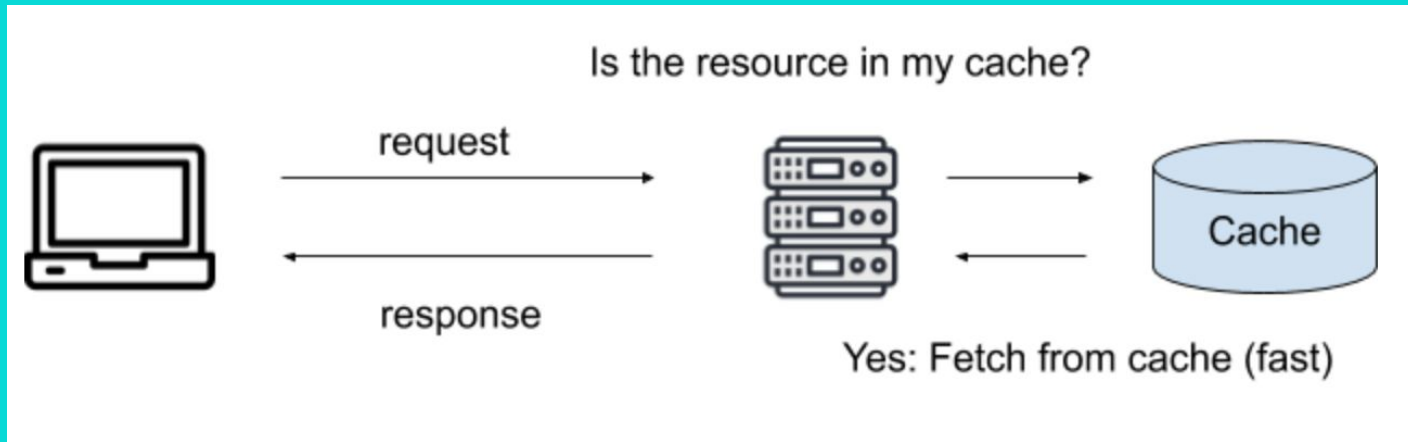
Client-side caching



Still fails at Phase 3



Server-side Caching



Server-side caching

```
<% cache(Event.cache_key_for_eventList) do %>
  ...
  <% @user.events.each do |event| %>
    <% cache(Event.cache_key_for_event(event))
do %>
  ...
  <% end %>
<% end %>
<% end %>
```

_eventcart.html.erb

```
def Event.cache_key_for_event(e)
  "event-#{e.id}-#{e.updated_at}-#{e.members.count}"
end

def Event.cache_key_for_eventList
  "eventList-#{Event.maximum(:updated_at)}-#{Member.maximum(:updated_at)}"
end
```

event.rb

Server-side caching

```
<div class="LI_GamerCardContainer">
  ...
  <% @userlist.each do |user| %>
    <% cache(User.cache_key_for_user(x)) do %>
      ...
    <% end %>
  <% end %>
</div>
```

```
def User.cache_key_for_user(x)
  "user-#{x.id}-#{x.updated_at}-#{x.events.count}"
end
```

Users

Server-side caching

```
<div class="columns" >
  <div class="column is-3" style="overflow: auto;">
    <% @conversations.each do |conversation| %>
      <% cache(Conversation.cache_key_for_conversation(conversation)) do %>
        ...
      <% end %>
    <% end %>
    ...
  <div style="height: 90%; overflow-y: scroll" id="message-window">
    <% @messages.reverse.each do |message| %>
      <% cache(Message.cache_key_for_message(message)) do %>
        ...
      <% end %>
    <% end %>
  </div>
  ...
</div>
```

Conversations


```
def
  Conversation.cache_key_
    for_conversation(x)

  "conversation-#{x.send_
    id}-#{x.recv_id}-#{x.up
    dated_at}"
end
```


```
def
  Message.cache_key_for_m
    essage(x)

  "message-#{x.user_id}-#
    {x.conversation_id}-#{x
    .created_at}"
end
```


Server-side caching



Fails at Phase 4 (originally Phase 3)



Read-slave



Read-slave

```
# config/database.yml
...

production:
  primary:
    adapter: postgresql
    database: <%= ENV['RDS_DB_NAME'] %>
    encoding: UTF8
    host: <%= ENV['RDS_HOSTNAME'] %>
    password: <%= ENV['RDS_PASSWORD'] %>
    port: <%= ENV['RDS_PORT'] %>
    username: <%= ENV['RDS_USERNAME'] %>
  primary_replica:
    adapter: postgresql
    database: <%= ENV['RDS_DB_NAME'] %>
    encoding: UTF8
    host: ""
    password: <%= ENV['RDS_PASSWORD'] %>
    port: <%= ENV['RDS_PORT'] %>
    username: <%= ENV['RDS_USERNAME'] %>
    replica: true
```

Read-slave

```
# app/models/application_record.rb
class ApplicationRecord < ActiveRecord::Base
  self.abstract_class = true
  connects_to database: { writing: :primary, reading: :primary_replica }
end
```

Read-slave

```
# config/application.rb
module Linksin
  class Application < Rails::Application
    ...

    config.active_record.database_selector =
      { delay: 2.seconds }

    config.active_record.database_resolver =
      ActiveRecord::Middleware::DatabaseSelector::Resolver

    config.active_record.database_resolver_context =
      ActiveRecord::Middleware::DatabaseSelector::Resolver::Session


  end
end
```

Read-slave


```
# config/database.yml
...

production:
  ...
  primary_replica:
    ...
    host: "linksin.civkgqjsrtyg.us-west-2.rds.amazonaws.com"
    ...
```

Read-slave



w/o read-slave




w/ read-slave

Read-slave

Instance	Max Phase Handled
w/o read-slave	5 (6 users/sec)
w/ read-slave	6 (10 users/sec)



SQL Optimization



SQL Optimization

- Used eager loading to circumvent N+1 query problem

```

rendering conversations/matches.html.erb within layouts/application
Converting Load (0.2ms) SELECT "conversations".* FROM "conversations" WHERE ((conversations.send_id = 1 OR conversations.recv_id = 1))
↳ app/views/conversations/matches.html.erb:3
app/view/conversations/matches.html.erb:5
Profile load (0.1ms) SELECT "profiles".* FROM "profiles" WHERE "profiles"."user_id" = ? LIMIT ? [{"user_id", 2}, {"LIMIT", 1}]
↳ app/views/conversations/matches.html.erb:20
app/view/conversations/matches.html.erb:5
app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 1}, {"LIMIT", 1}]
↳ app/views/conversations/matches.html.erb:6
app/view/conversations/matches.html.erb:6
Profile load (0.1ms) SELECT "profiles".* FROM "profiles" WHERE "profiles"."user_id" = ? LIMIT ? [{"user_id", 4}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:20
CACHE User load (0.0ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 1}, {"LIMIT", 1}]
↳ app/views/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 11}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:6
Profile load (0.1ms) SELECT "profiles".* FROM "profiles" WHERE "profiles"."user_id" = ? LIMIT ? [{"user_id", 11}, {"LIMIT", 1}]
↳ app/views/conversations/matches.html.erb:20
CACHE User load (0.0ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 1}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:5
app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 6}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:6
Profile load (0.2ms) SELECT "profiles".* FROM "profiles" WHERE "profiles"."user_id" = ? LIMIT ? [{"user_id", 6}, {"LIMIT", 1}]
↳ app/views/conversations/matches.html.erb:20
app/view/conversations/matches.html.erb:5
app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 8}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:6
app/view/conversations/matches.html.erb:6
app/view/conversations/matches.html.erb:5
app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 12}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:20
CACHE User load (0.0ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 1}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 12}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:6
Profile load (0.1ms) SELECT "profiles".* FROM "profiles" WHERE "profiles"."user_id" = ? LIMIT ? [{"user_id", 12}, {"LIMIT", 1}]
↳ app/views/conversations/matches.html.erb:20
CACHE User load (0.0ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 1}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 15}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:6
Profile load (0.1ms) SELECT "profiles".* FROM "profiles" WHERE "profiles"."user_id" = ? LIMIT ? [{"user_id", 15}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:20
CACHE User load (0.0ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 1}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"id", 5}, {"LIMIT", 1}]
↳ app/view/conversations/matches.html.erb:6
app/view/conversations/matches.html.erb:6
app/view/conversations/matches.html.erb:5
app/view/conversations/matches.html.erb:5
User load (0.1ms) SELECT "users".* FROM "users" WHERE "users"."id" = ? LIMIT ? [{"user_id", 5}, {"LIMIT", 1}]

```

Eager Loading

[illegible]

SQL Optimization

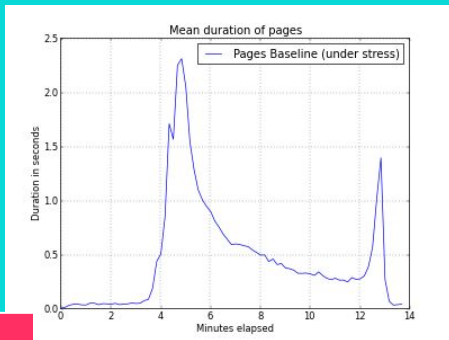
- Attempted (keyword: attempted) to speed up user matching with a custom query

```
def show
  @user = User.find(params[:id])
  @userlist = randomShow(@user).select do
    |user|
      conversation =
        Conversation.between(user.id, @user.id)
      conversation.empty?
    end
  end
end
```

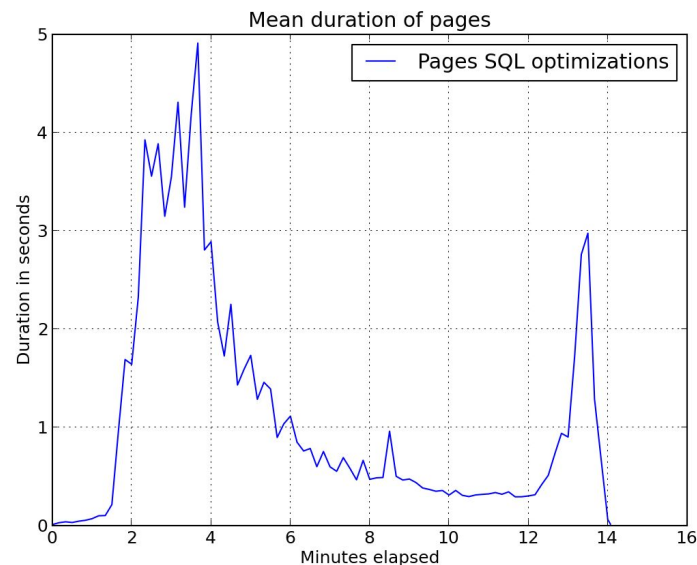
```
def randomShow(user)
  User.find_by_sql("SELECT * FROM users u
    WHERE ((u.id != #{user.id}) AND
    NOT EXISTS
    ( SELECT u.id
      FROM conversations c
      WHERE (c.send_id = u.id) OR
      (c.recv_id = u.id)
    ))
    ORDER BY RANDOM()
    LIMIT 20;")
end
```

SQL Optimization

- Performance actually decreased!
 - NOT EXISTS runs the sub-query for each id → scales horribly
 - NOT IN should have been used instead as the sub-queries are only run once



Before



After

SQL Optimization

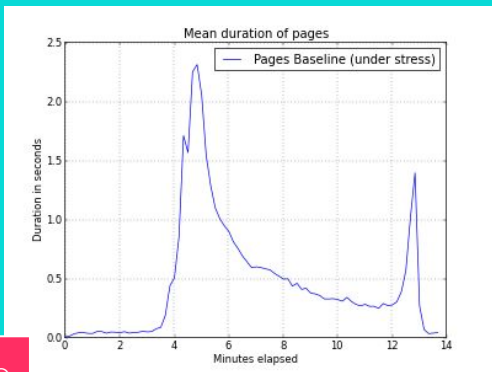
Instance	Max Phase Handled
w/o sql changes	3 (2 users/sec)
w/ sql changes	1 (1 users/sec)




Pagination

Pagination

- Switched from will_paginate to pagy for countless pagination
- Used pagination on user for matching → load more users via ajax when down to two cards



Before



After

Pagination

Instance	Max Phase Handled
w/o pagination	3 (2 users/sec)
w/ pagination	4 (4 users/sec)

Future Development

Database sharding, more features



- ❑ Database Sharding
- ❑ Improve Events
- ❑ Real-time messaging
- ❑ Notifications
- ❑ Linking external accounts
- ❑ Combine optimizations

Conclusions & Lessons Learned




Agile Development



Pair Programming



Scalable Web Service



Load Testing



Thanks!